


### Assignment 1: Feedforward neural networks (approximate training run-time on a CPU: 5 min)

Implement a feedforward neural network that solves the speech classification task. This neural network must have four hidden layers with 128 neurons each. Apply a rectified linear unit (ReLU) activation function at the output of each hidden layer.



In addition, consider a minibatch size of 1,024 samples and, as the optimizer, employ stochastic gradient descent with a learning rate of 0.01, momentum of 0.9 and weight decay of 0.00001. Train your model for 50 epochs. Use “torch.manual\_seed(0)” to make your results reproducible.

- 1) What is the size of your input layer and why? What output layer activation function and training loss function have you chosen and why?
- 2) Plot the training and validation losses, as well as the training and validation accuracies, as a function of the training iteration<sup>1</sup>. Do you think that overfitting has occurred? If your answer was yes, how may you solve it? (Just explain with words)
- 3) What is the test accuracy?
- 4) Calculate, by hand, the total number of parameters of the model, and indicate, step by step, how you reached your solution.

#### Hints:

- Do not forget to normalize your speech feature matrices, e.g.,  $X_{train} \leftarrow (X_{train} - \text{np.mean}(X_{train})) / \text{np.std}(X_{train})$
- Shuffle your training data before model training.
- Accuracy is defined as the ratio between the number of correct classifications and the total number of classifications.
- To check that you correctly calculated by hand the number of parameters of your model, you may want to use the following snip of code:

```
def get_no_params(model):
    nop = 0
    for param in list(model.parameters()):
        nn = 1
        for s in list(param.size()):
            nn = nn * s
        nop += nn
    return nop
```

---

<sup>1</sup> Note that your training data set is comprised of [9,489 training samples / 1,024 samples per minibatch] = 10 minibatches. Hence, your training algorithm runs for 50 epochs × 10 minibatches per epoch = 500 iterations.