

### Assignment 3: Recurrent neural networks (approximate training run-time on a CPU: 1 hour)

Implement a recurrent neural network that solves the speech classification task. This neural network must have two gated recurrent unit (GRU) layers with 64-dimensional hidden states followed by two fully-connected layers with 128 neurons each. Apply a rectified linear unit (ReLU) activation function at the output of each fully-connected layer. Furthermore, bear in mind that each speech feature matrix represents a temporal sequence of  $T=101$  frames. Hence, for classification, you only have to pass the hidden state of the second GRU layer at time  $T$  to the first fully-connected layer.



Once again, consider a minibatch size of 1,024 samples and, as the optimizer, employ Adam with default parameters. Train your model for 50 epochs. Use “torch.manual\_seed(0)” to make your results reproducible and be patient... backpropagation through time (BPTT) can be computationally expensive!

- 1) Plot the training and validation losses, as well as the training and validation accuracies, as a function of the training iteration.
- 2) What is the test accuracy?
- 3) Calculate, by hand, the total number of parameters of the model, and indicate, step by step, how you reached your solution.
- 4) Very briefly, compare, in terms of performance and computational complexity, this model with the models that you implemented in the previous assignments.

#### Hints:

- Do not forget to normalize your speech feature matrices, e.g.,  $X_{train} \leftarrow (X_{train} - \text{np.mean}(X_{train})) / \text{np.std}(X_{train})$
- Shuffle your training data before model training.
- Accuracy is defined as the ratio between the number of correct classifications and the total number of classifications.
- Training on a CPU takes time. To check more quickly that your implementation is correct, you may want to train your model for 5 epochs. In this case, test accuracy should be *around* 64%.
- To check that you correctly calculated by hand the number of parameters of your model, you may want to use the following snip of code:

```
def get_no_params(model):
    nop = 0
    for param in list(model.parameters()):
        nn = 1
        for s in list(param.size()):
            nn = nn * s
        nop += nn
    return nop
```